



Embedded and Real-Time Systems

Spring 2021

Hamed Farbeh

farbeh@aut.ac.ir

Department of Computer Engineering

Amirkabir University of Technology

Lecture 11

Copyright Notice

This lecture is adopted from
IN4343 Real-Time Systems Course 2018 – 2019, Mitra Nasri,
Delft University of Technology

Non-preemptive Scheduling

Paper 1:

Mitra Nasri and Gerhard Fohler, "Non-Work-Conserving Non-preemptive Scheduling: Motivations, Challenges, and Potential Solutions," in the Proceedings of the Euromicro Conference on Real-Time Systems (ECRTS), 2016, pp. 165-175.


Paper 2:

Mitra Nasri and Björn B. Brandenburg, "An Exact and Sustainable Analysis of Non-Preemptive Scheduling", in the Proceedings of the Real-Time Systems Symposium (RTSS), 2017, pp. 1-12

Disclaimer: A few slides have been taken from [Giorgio Buttazzo's](http://www.giorgio Buttazzo.it) website:

<http://retis.sssup.it/~giorgio/rtS-MECS.html>

Buttazzo's book, chapter 4, 8



Agenda

- EDF schedulability analysis

Book: chapter 4

- Optimality of EDF for periodic tasks
- EDF response-time analysis

- Non-preemptive scheduling

- Why non-preemptive execution?
- Existing schedulability analysis for NP-FP and NP-EDF
- Behind the scenes!

Book: chapter 8

Book: chapter 8

Paper 1, 2

Things you need to know about work-conserving NP scheduling!



- What is the hyperperiod of
5, 10, 20?

20

- What is the hyperperiod of
5, 10, 21?

210

Optimality

The proof comes from Dertouzos 1974

Schedulability test

- Liu and Layland 1973

A task set τ is schedulable by EDF **if and only if:**

$$\sum_{i=1}^n \frac{C_i}{T_i} \leq 1$$

Assumptions:

- Fully preemptive
- Independent tasks
- No self-suspension
- No scheduling overhead
- $\forall i, D_i = T_i$

EDF schedulability analysis for $D < T$

Processor Demand Criterion [Baruah '90]

In **any interval** of length L , the **computational demand** $g(0, L)$ of the task set must be **no greater** than L .


$$\forall L > 0, \quad g(0, L) \leq L$$

Applicable to periodic tasks with $\forall i, \phi_i = 0$


Understanding the processor demand (g function)

- The demand in $[t_1, t_2]$ is the computation time of those tasks arrived at or after t_1 with deadline less than or equal to t_2 :

$$g(t_1, t_2) = \sum_{\forall J_{k,j}, t_1 \leq a_{k,j} < d_{k,j} \leq t_2} C_k$$



Demand of a periodic task set



$$g_i(0, L) = \max \left\{ 0, \left\lfloor \frac{L - D_i + T_i}{T_i} \right\rfloor \cdot C_i \right\}$$


$$g_i(0, L) = \max \left\{ 0, \left\lfloor \frac{L - D_i}{T_i} + 1 \right\rfloor \cdot C_i \right\}$$

$$g_i(0, L) = \max \left\{ 0, \left(\left\lfloor \frac{L - D_i}{T_i} \right\rfloor + 1 \right) \cdot C_i \right\}$$

$$g(0, L) = \sum_{i=1}^n g_i(0, L)$$


Examples

$$g_i(0, L) = \max \left\{ 0, \left\lfloor \frac{L - D_i}{T_i} + 1 \right\rfloor \cdot C_i \right\}$$



L	$g_2(0, L)$	L	$g_2(0, L)$
4	$g_2(0, 4) = \left\lfloor \frac{4 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$	9	$g_2(0, 9) = \left\lfloor \frac{9 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$
5	$g_2(0, 5) = \left\lfloor \frac{5 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$	10	$g_2(0, 10) = \left\lfloor \frac{10 - 4}{6} + 1 \right\rfloor \cdot 2 = 4$
6	$g_2(0, 6) = \left\lfloor \frac{6 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$	11	$g_2(0, 11) = \left\lfloor \frac{11 - 4}{6} + 1 \right\rfloor \cdot 2 = 4$
7	$g_2(0, 7) = \left\lfloor \frac{7 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$	15	$g_2(0, 15) = \left\lfloor \frac{15 - 4}{6} + 1 \right\rfloor \cdot 2 = 4$
8	$g_2(0, 8) = \left\lfloor \frac{8 - 4}{6} + 1 \right\rfloor \cdot 2 = 2$	16	$g_2(0, 16) = \left\lfloor \frac{16 - 4}{6} + 1 \right\rfloor \cdot 2 = 6$

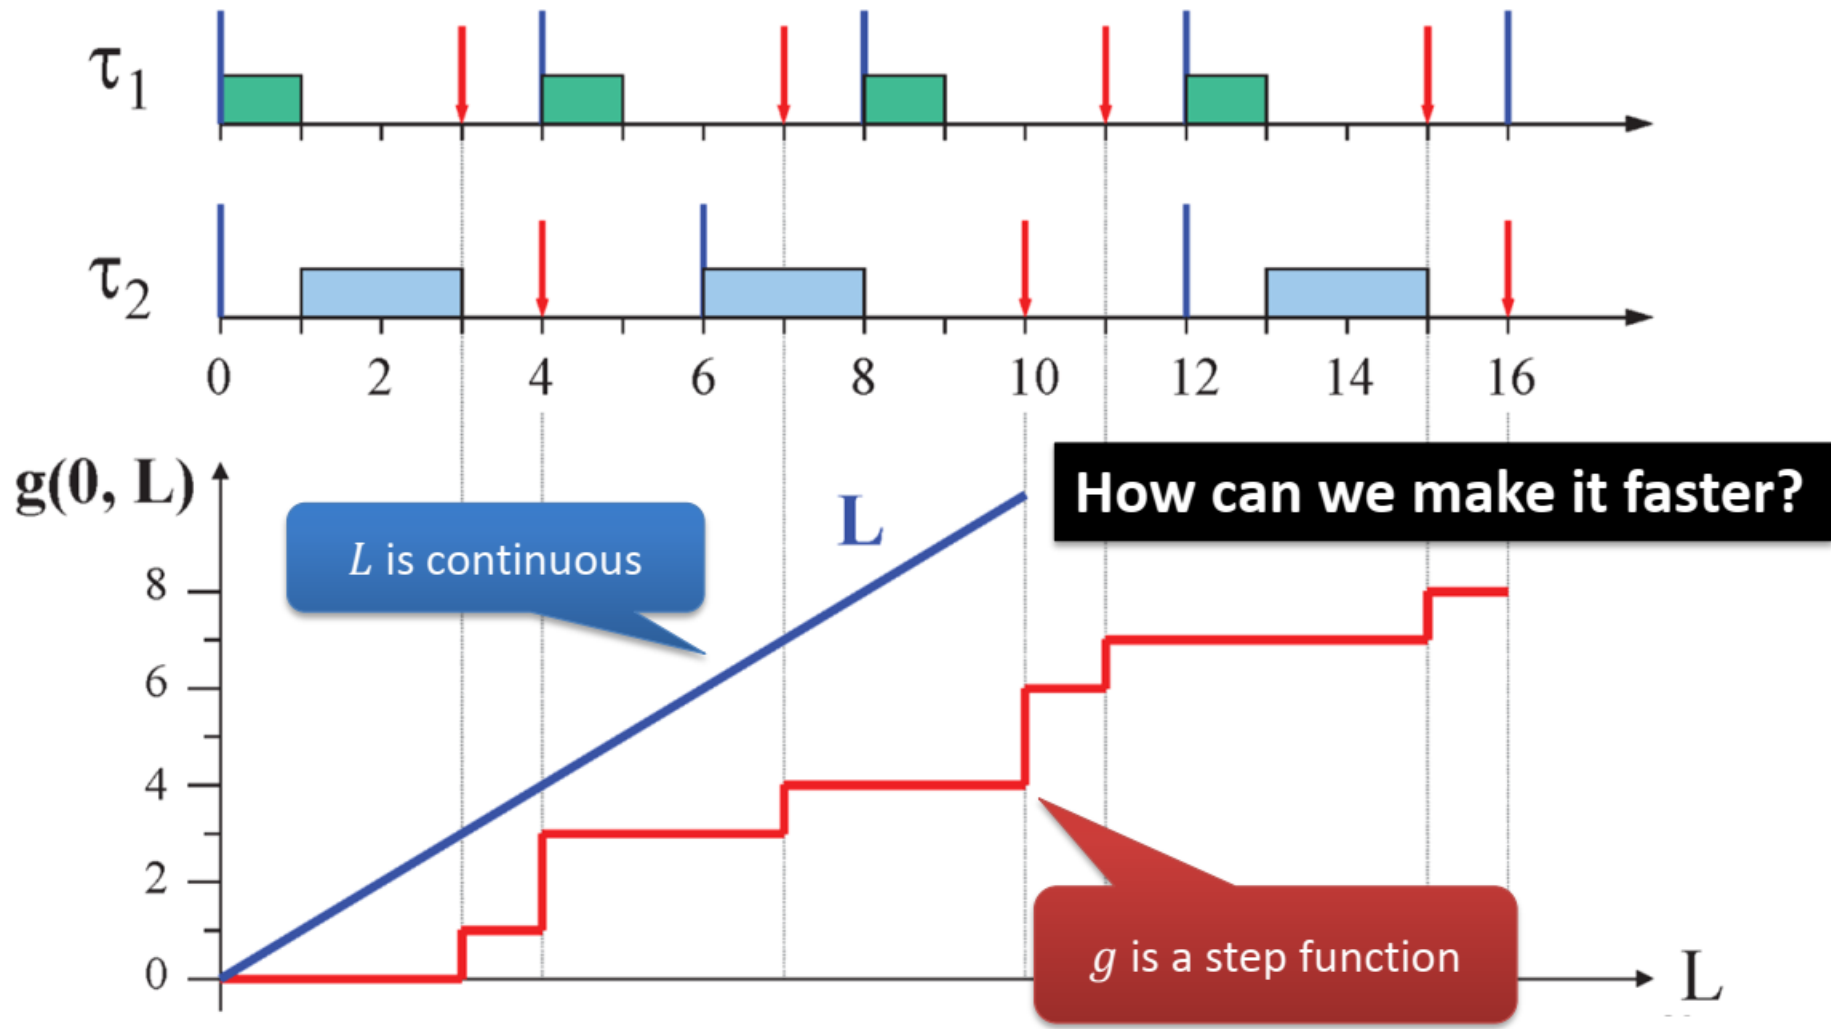
When does Demand Bound Function change?



$$\forall L > 0, \\ g(0, L) \leq L$$

$$g(0, L) = \sum_{i=1}^n g_i(0, L)$$

$$g_i(0, L) = \max \left\{ 0, \left\lfloor \frac{L - D_i}{T_i} + 1 \right\rfloor \cdot C_i \right\}$$



Bounding complexity

- Since $g(0, L)$ is a step function, we can check feasibility only at **deadline points**, where g changes.
- If tasks are synchronous and $U \leq 1$, we can check the feasibility up to the hyperperiod H :

$$H = lcm(T_1, \dots, T_n)$$

Bounding complexity: finding a safe upper bound for the analysis window

An upper bound on g



$$g(0, L) = \sum_{i=1}^n \left\lfloor \frac{L + T_i - D_i}{T_i} \right\rfloor \cdot C_i \leq G(0, L) = \sum_{i=1}^n \left(\frac{L + T_i - D_i}{T_i} \right) \cdot C_i$$

$$\boxed{\lfloor x \rfloor \leq x}$$

$$G(0, L) = \sum_{i=1}^n L \frac{C_i}{T_i} + \sum_{i=1}^n (T_i - D_i) \cdot \frac{C_i}{T_i}$$

$$= L \cdot U + \sum_{i=1}^n (T_i - D_i) \cdot U_i$$


$$g(0, L) < G(0, L)$$

Limiting L

$$U \leq 1$$

$$G(0, L) = L \cdot U + \sum_{i=1}^n (T_i - D_i) \cdot U_i$$

$$L^* = \frac{\sum_{i=1}^n (T_i - D_i) \cdot U_i}{1 - U}$$



For any $L \geq L^*$, it is meaningless to check if $g(0, L) \leq L$ because it will obviously be smaller!
Hence, we can stop the search at L^* .

Processor Demand Test

$$\forall L \in D, \quad g(0, L) \leq L$$

Where D is the set of deadline points for which $g(0, L)$ must be calculated:

$$D = \{d_{i,j} \mid d_{i,j} \leq \min\{H, L^*\}\}$$

$$H = lcm(T_1, \dots, T_n)$$

$$L^* = \frac{\sum_{i=1}^n (T_i - D_i) \cdot U_i}{1 - U}$$

FP

v.s.

EDF




EDF v.s. FP

τ_i	C_i	T_i	D_i
τ_1	2	5	5
τ_2	4	10	8
τ_3	4	20	17

Context switches


FP



How many preemptions?

3 preemptions in jobs $J_{2,1}$, $J_{3,1}$, $J_{2,2}$

EDF



How many preemptions?

1 preemption in job $J_{3,1}$

In average, FP has more context switches


EDF v.s. FP

What do you think will happen for FP if $U > 1$?

FP under permanent overload

➡ starvation for low priority tasks

$$U = \frac{4}{8} + \frac{6}{12} + \frac{5}{20} = 1.25$$




- High priority tasks execute at the proper rate
- Low priority tasks are completely blocked

EDF v.s. FP

EDF under permanent overload

$$U = \frac{4}{8} + \frac{6}{12} + \frac{5}{20} = 1.25$$



- All tasks execute at a slower rate
- No task is blocked

EDF v.s. FP

Schedulability analysis

	$D_i = T_i$	$D_i \leq T_i$
RM	<p><i>Suff.: polynomial</i> $O(n)$</p> <p>LL: $\sum U_i \leq n(2^{1/n} - 1)$</p> <p>HB: $\prod(U_i + 1) \leq 2$</p> <p><i>Exact pseudo-polynomial</i> RTA</p>	<p><i>pseudo-polynomial</i> Response Time Analysis</p> <p>$\forall i \quad R_i \leq D_i$</p> $R_i = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i}{T_k} \right\rceil C_k$
EDF	<p><i>polynomial:</i> $O(n)$</p> $\sum U_i \leq 1$	<p><i>pseudo-polynomial</i> Processor Demand Analysis</p> $\forall L > 0, \quad g(0, L) \leq L$

FP's response-time analysis is usually faster than EDF

EDF v.s. FP

EDF

- Higher schedulability
- Reduces context switches
- During overloads does not starve low-priority tasks

FP

- Simpler to implement
- Widely implemented in most operating systems
- More predictable during overloads

How does DM compare to EDF?

Agenda

- EDF schedulability analysis
 - Optimality of EDF for periodic tasks
 - EDF response-time analysis
- Non-preemptive scheduling
 - **Why non-preemptive execution?**
 - **Existing schedulability analysis for NP-FP and NP-EDF**
 - **Behind the scenes!**


Things you need to know about work-conserving NP scheduling!



Disadvantages of preemptions

1- Context switch cost


It is the time taken by the scheduler to suspend the running task, switch the context, and dispatch the new incoming task.



Disadvantages of preemptions

2- Cache-related preemption delay (CRPD)

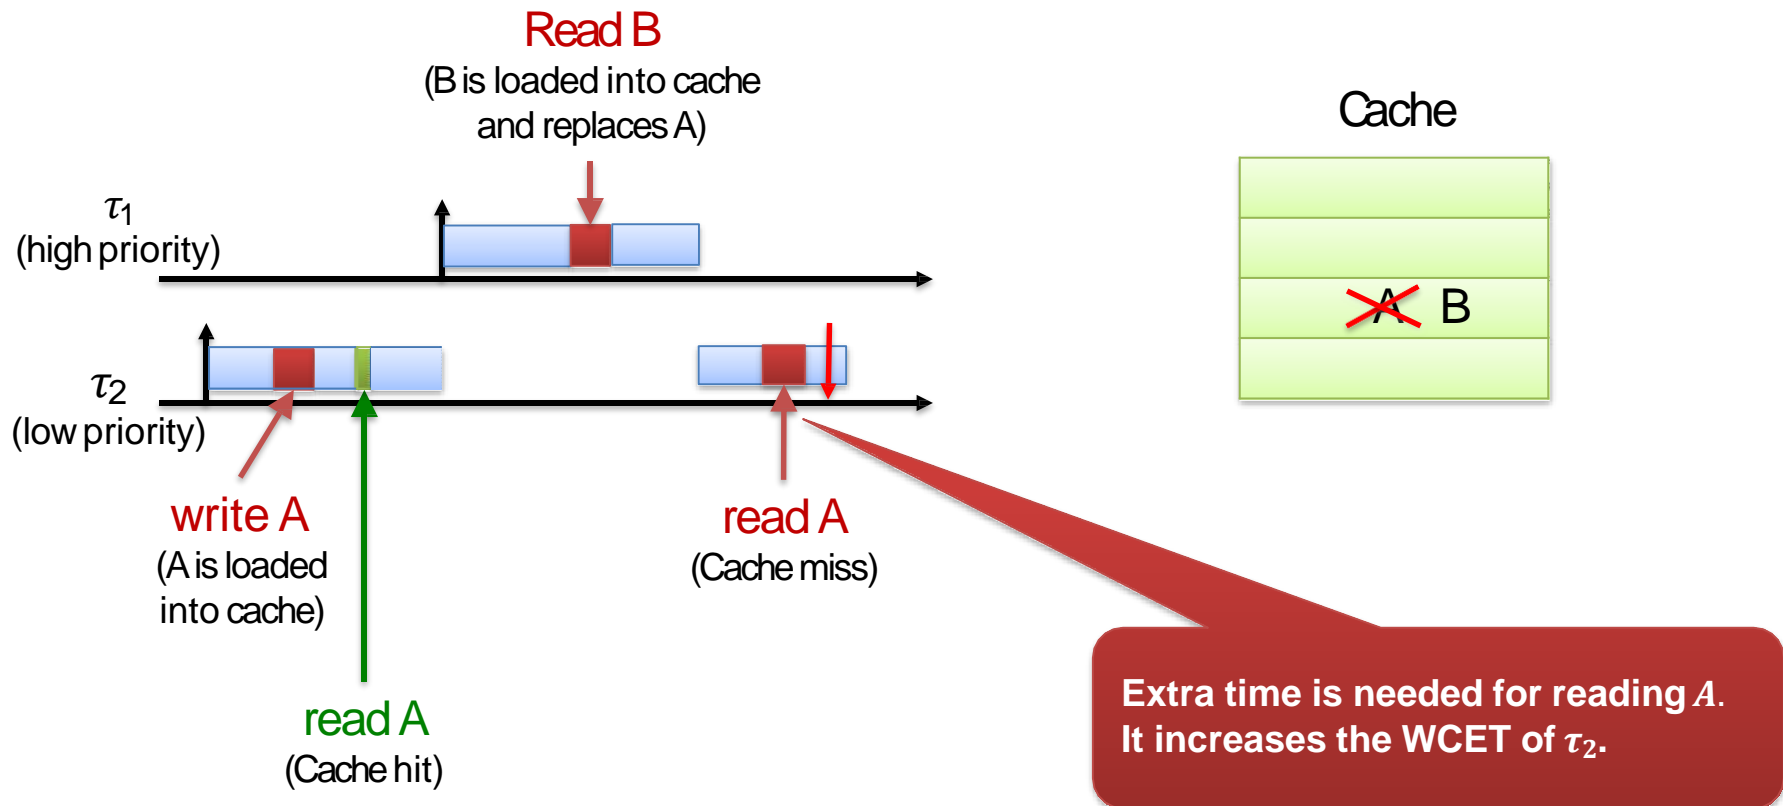
It is the delay introduced by high-priority tasks that **evict cache lines** containing data used in the future:



Disadvantages of preemptions

2- Cache-related preemption delay (CRPD)

It is the delay introduced by high-priority tasks that **evict cache lines** containing data used in the future:




Disadvantages of preemptions

3- larger worst-case execution time

Preemptions cause CRPD which in turn increases the WCET of the task

The amount of CRPD depends on the number of preemptions a job suffers

Task experiencing preemptions by higher priority tasks:



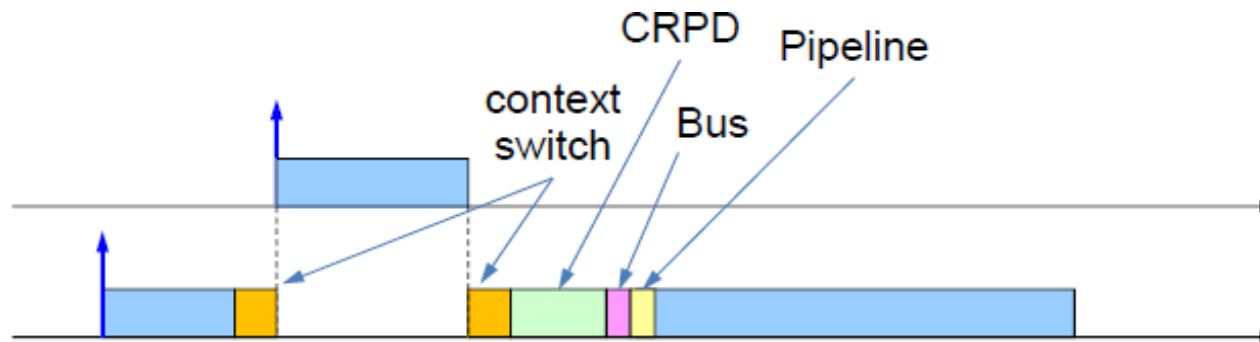
Disadvantages of preemptions

4 Pipeline cost

time to flush the pipeline when a task is interrupted and to refill it when task is resumed.

5 Bus cost

time spent waiting for the bus due to additional conflicts with I/O devices, caused by extra accesses to the RAM for the extra cache misses.




Disadvantages of preemptions


A preemption may cause more preemptions

the extra execution time also increases the number of preemptions:

Preemption overhead is zero




Preemption overhead is NOT zero




Disadvantages of preemptions

Preemption cost can be very large

- WCETs may increase up to 35% in the presence of preemptions (less efficiency)!



In preemptive execution, the execution times are more variable (**less predictability**)

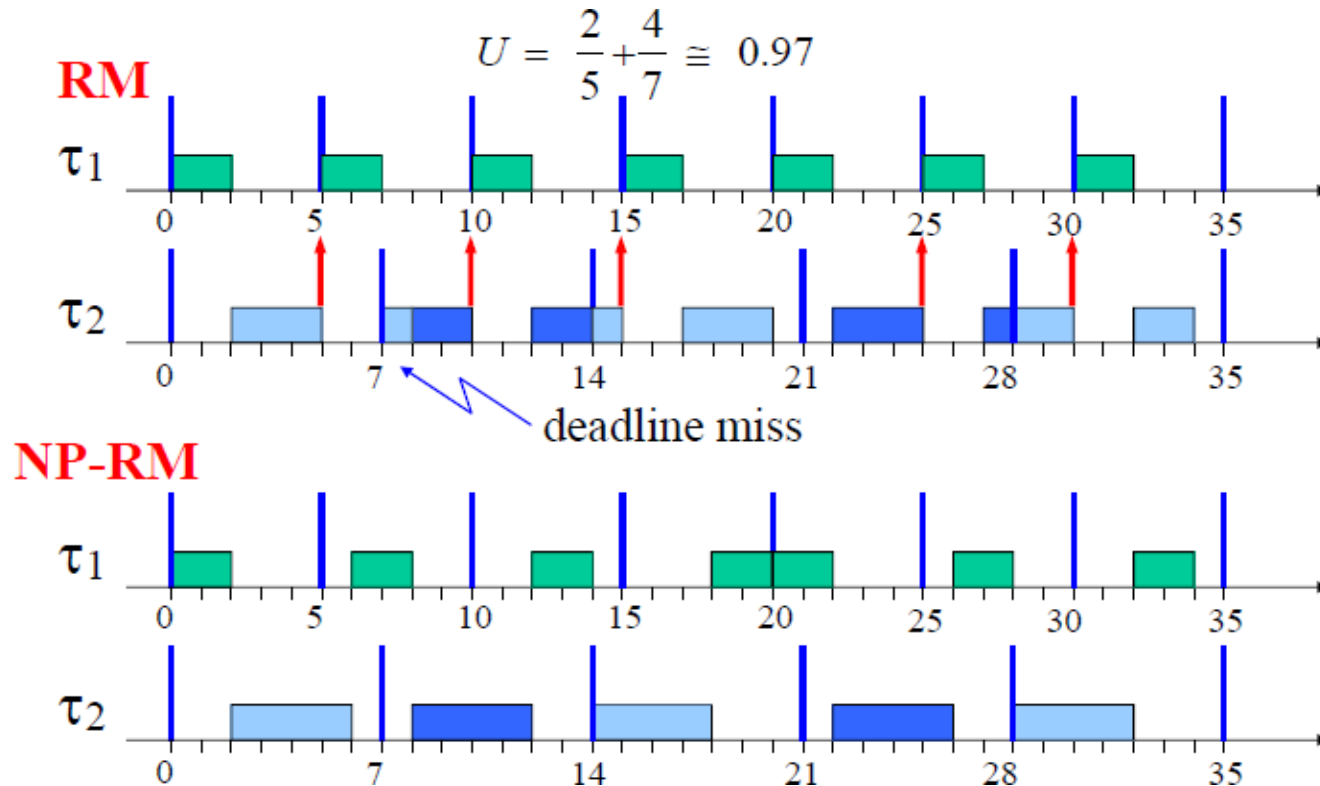


Advantages of NP scheduling

- It reduces **context-switch overhead**:
 - Making WCETs smaller and more predictable
- It simplifies the **access to shared resources**:
 - No semaphores are needed for critical sections
- It reduces **stack size**:
 - Task can share the same stack, since no more than one task can be in execution
- It allows achieving **small I/O Jitter**:
 - “finishing_time – start_time” has a low variation

Advantages of NP scheduling

- In fixed-priority systems, non-preemptive execution can even improve schedulability (in some cases)




**So, why
preemptive execution?**




Disadvantages of NP scheduling

- In general, NP scheduling reduces schedulability because of introducing **blocking delays** on the high-priority tasks



Disadvantages of NP scheduling

- The utilization bound under non-preemptive scheduling drops to zero




$$U_1 + U_2 \sim 0$$

Infeasible! Despite having $U \sim 0$

Disadvantages of NP scheduling

Anomalies



Anomaly is a situation in which a deadline miss happens when we don't expect it to happen!

Example: the task set is feasible on the current processor but when we use a faster one, it becomes unschedulable!